

EQUITY	COMPETITION P&L	TRADES	WIN RATE	MAX DRAWDOWN
\$99,986	-\$14 (-0.01%)	9 closed	67%	~\$139

Snapshot as of Sep 3, 4:00 PM ET (Thu close) – the agent continues trading autonomously through the submission deadline. Authoritative P&L: paper account PA3KGE66JM72 · live view on the dashboard.

STRATEGY

A two-sleeve, defined-risk options book on Alpaca paper trading. The core sleeve sells short-dated (0–5 DTE) credit spreads on XSP/SPY/QQQ at ~15-delta short strikes (~1σ) — harvesting the variance risk premium, the one edge that reliably realizes inside a single week because theta accrues every session. The catalyst sleeve expresses event-driven views as small debit spreads. Every structure has defined max loss at entry; exits are rule-based (60% profit-take, 2×-credit stop, expiry handling). An expiry ladder refuses thin credit — a missing 15Δ trade becomes no trade, never a 30Δ trade — and all positions expire by Sep 4 so P&L; is realized inside the judging window.

AI LOGIC – THE LLM PROPOSES, DETERMINISTIC CODE DISPOSES

RESEARCH	STRATEGY	RISK (veto power)	EXECUTION
----------	----------	-------------------	-----------

Each morning the research agent — Claude, via the Claude Agent SDK — connects to Alpaca's official MCP server and reads the market through its tools (snapshots, news, movers, option chains; order-placing tools are whitelisted out and explicitly banned), plus web search for the macro calendar. It emits a structured brief that sets the day's posture: delta cap, size multiplier, entry blackouts, stand-down. Its influence is tighten-only, enforced in code — every recommendation is min()-ed against config, so the LLM can make the system more careful but structurally cannot make it more aggressive. It earned its keep on day one: it caught that our static calendar had the wrong kickoff-day event (the new Fed Chair's Jackson Hole keynote) and set the entry blackout accordingly. Every proposal, veto, fill, and brief is journaled with written rationale, feeding Telegram reports and the live dashboard: the agent explains every decision it ever made.

RISK GATES – DETERMINISTIC, NO LLM IN THE LOOP

1 · Daily-loss kill switch	–2% (\$2,000)	blocks new risk; managing/closing always allowed
2 · Drawdown kill switch	–5% (\$5,000)	flatten everything + freeze for the competition
3 · Per-trade max loss	\$1,500	shrink-before-veto sizing
4 · Open-risk budget	\$10,000	sum of defined max losses across the book
5 · Position count / 6 · Order size	8 positions · 10 contracts	hard caps
7 · Weekend rule	Friday = ODTE only	nothing survives a weekend
8 · Market hours	15m open / 30m close buffers	no entries in opening chaos or gamma hour

Kill switches fire only on sustained, sanity-checked breaches — N consecutive valid readings over minutes. Readings failing data-quality checks (crossed/stale/thin quotes, implausible jumps, broker-vs-computed equity disagreement) pause the clock rather than trigger action: a data glitch can cause nothing, never a liquidation. Exits are always limit orders re-pegged toward the market — never market orders into an options book. Live validation: a position that dipped to –\$94 intraday recovered fully because the debounced stop refused to panic.

ALPACA INFRASTRUCTURE

Trading API	Multi-leg MLEG orders for atomic spread execution — signed net limit price, *_TO_CLOSE intents on exits, buying-power reduction verified live to equal defined max loss. Positions, clock/calendar, portfolio history.
MCP server	The research agent's entire market-reading surface (official alpaca-mcp-server) — core of the AI logic above.
CLI	Independent account cross-check in every end-of-day report — two code paths must agree on equity (official alpacahq/cli).
Paper env	Two full pre-kickoff shakedown sessions verified every path live: fills, re-pegs, profit-takes, kill-switch plumbing, restart safety (startup sweep of orphaned orders). Hard kickoff gate guaranteed zero order activity before the competition.